Alex Chen

Senior AI Engineer | Generative AI & LLM Systems

alex.chen@cynbit-demo.com | +1 (555) 234-8901 | San Francisco, CA | github.com/alexchen-ai

PROFESSIONAL SUMMARY

Results-driven AI Engineer with 4+ years of experience designing and deploying scalable deep learning architectures, Retrieval-Augmented Generation (RAG) pipelines, and autonomous LLM agents in production environments. Proficient in PyTorch, FastAPI, Vector Databases, and Dockerized microservices.

TECHNICAL SKILLS

Python • PyTorch • FastAPI • Docker • RAG • LangChain • Vector Databases (ChromaDB, Pinecone) • Git • Hugging Face • Natural Language Processing (NLP) • MLflow • REST APIs • PostgreSQL

PROFESSIONAL EXPERIENCE

AI Systems Engineer | NeuralScale Tech *(2023 - Present)*

- Architected enterprise RAG system with LangChain and Pinecone vector databases, cutting support ticket resolution time by 48%.
- Built asynchronous microservices in FastAPI and Docker serving PyTorch neural embeddings with <45ms p99 latency.
- Implemented CI/CD pipelines with Git and Docker for automated model benchmarking and zero-downtime deployment.

Machine Learning Associate | Cognitive Cloud Labs *(2021 - 2023)*

- Trained and fine-tuned Transformer-based NLP models using Hugging Face and PyTorch on multi-GPU nodes.
- Integrated MLflow for experiment tracking, hyperparameter optimization, and model artifact governance.
- Developed Python ETL scripts processing 15M+ text documents for downstream retrieval tasks.

EDUCATION

B.S. in Computer Science — University of California, Berkeley (2017 - 2021)